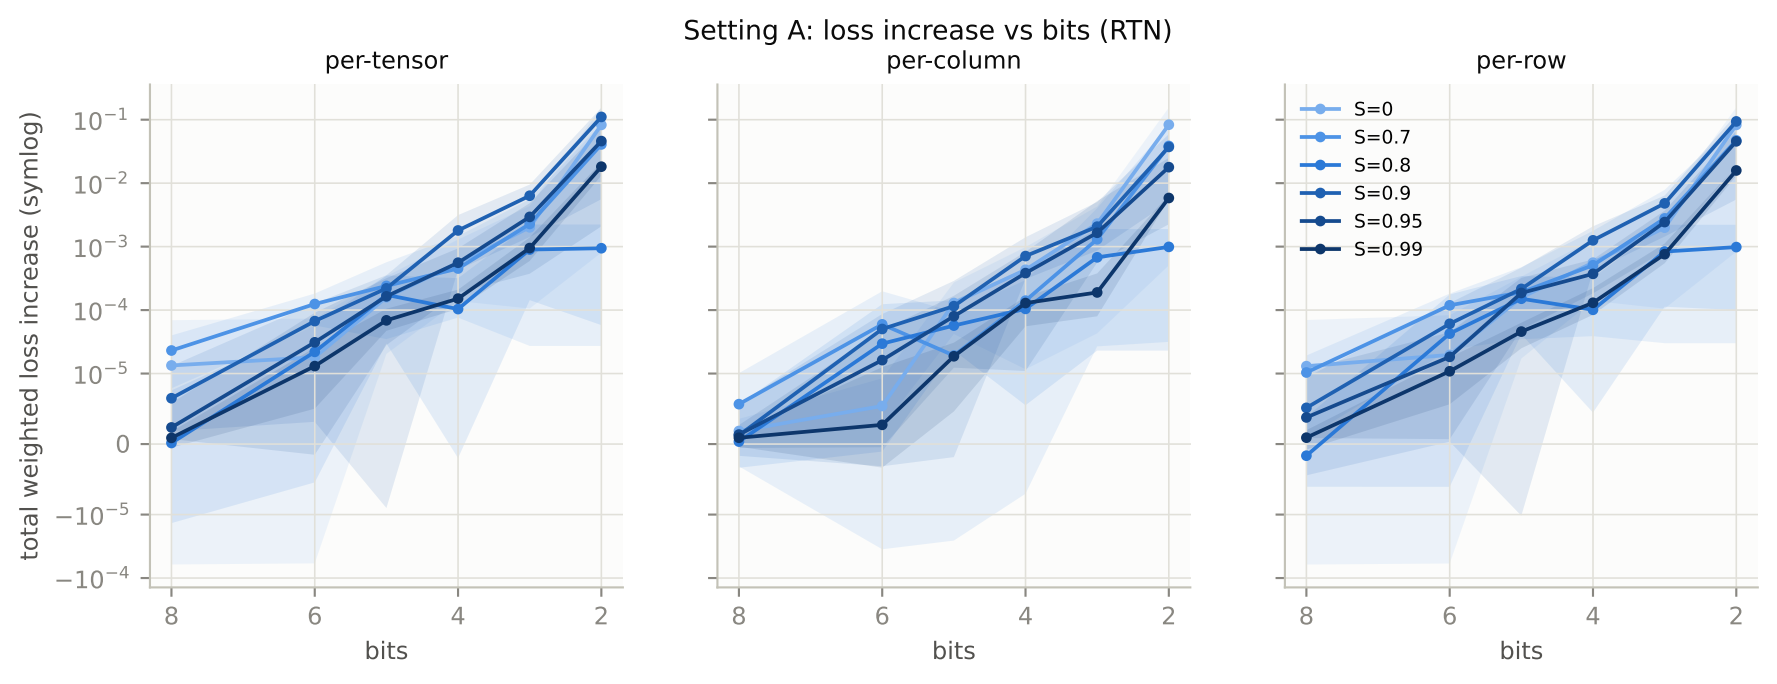
Setting A: loss increase vs bits (RTN)
per-tensor
per-column
per-row
total weighted loss increase (symlog)
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
0
-10^{-5}
-10^{-4}
8
6
4
2
bits
S=0
S=0.7
S=0.8
S=0.9
S=0.95
S=0.99