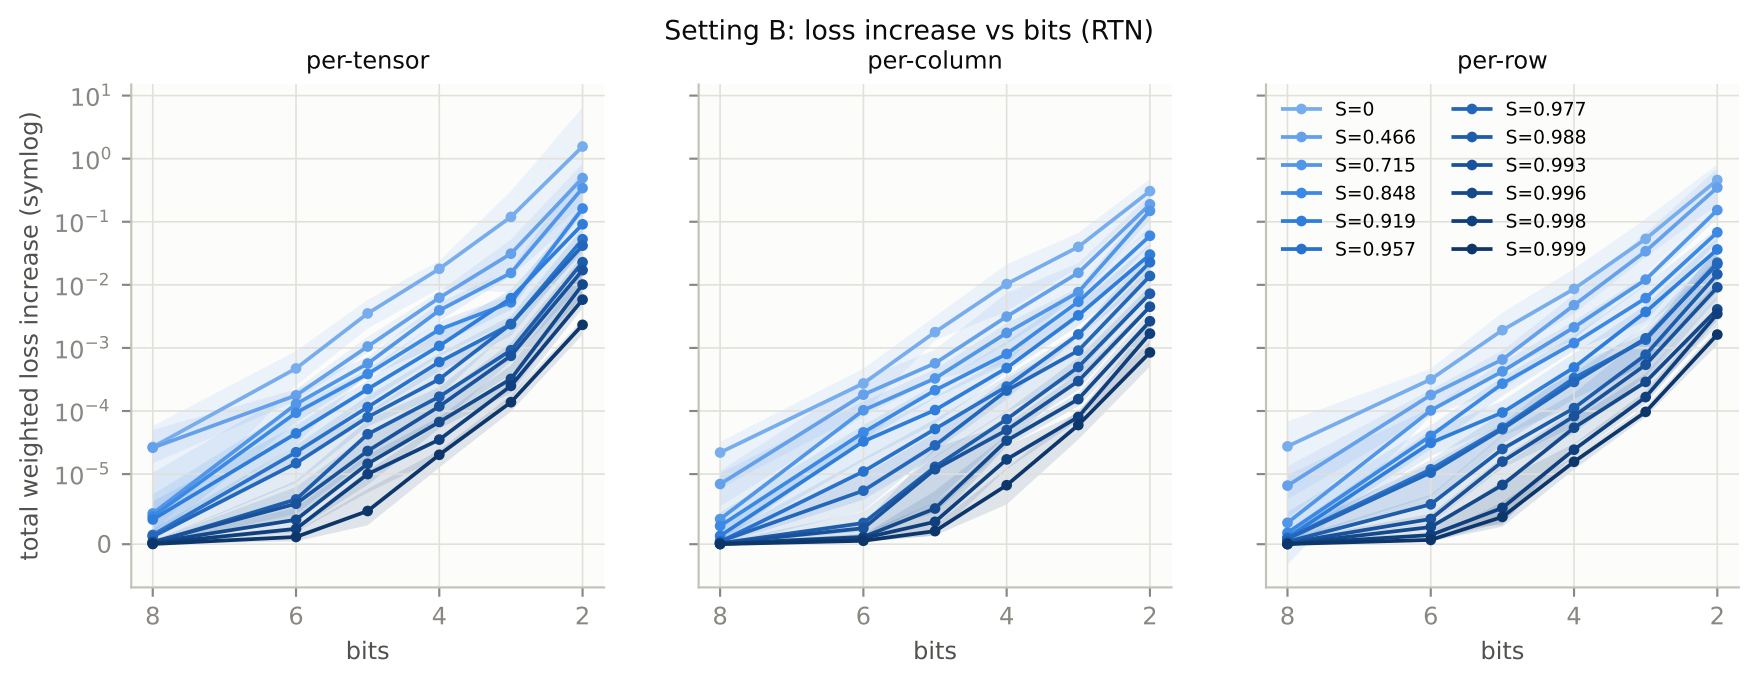
Setting B: loss increase vs bits (RTN)
per-tensor
per-column
per-row
total weighted loss increase (symlog)
10^{1}
10^{0}
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
0
8
6
4
2
bits
S=0
S=0.466
S=0.715
S=0.848
S=0.919
S=0.957
S=0.977
S=0.988
S=0.993
S=0.996
S=0.998
S=0.999