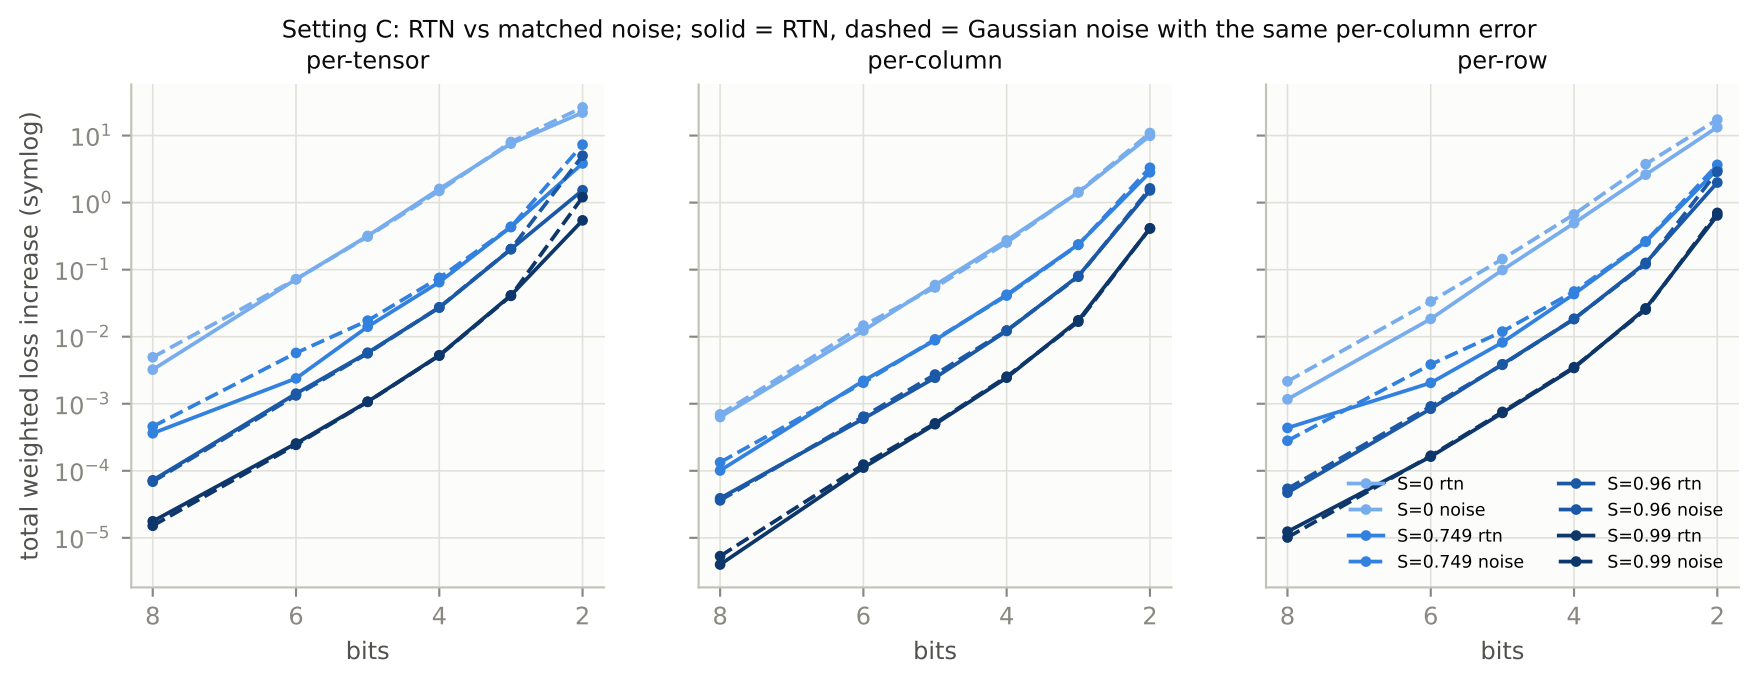
Setting C: RTN vs matched noise; solid = RTN, dashed = Gaussian noise with the same per-column error
per-tensor
per-column
per-row
total weighted loss increase (symlog)
10^{1}
10^{0}
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
8
6
4
2
bits
S=0 rtn
S=0 noise
S=0.749 rtn
S=0.749 noise
S=0.96 rtn
S=0.96 noise
S=0.99 rtn
S=0.99 noise