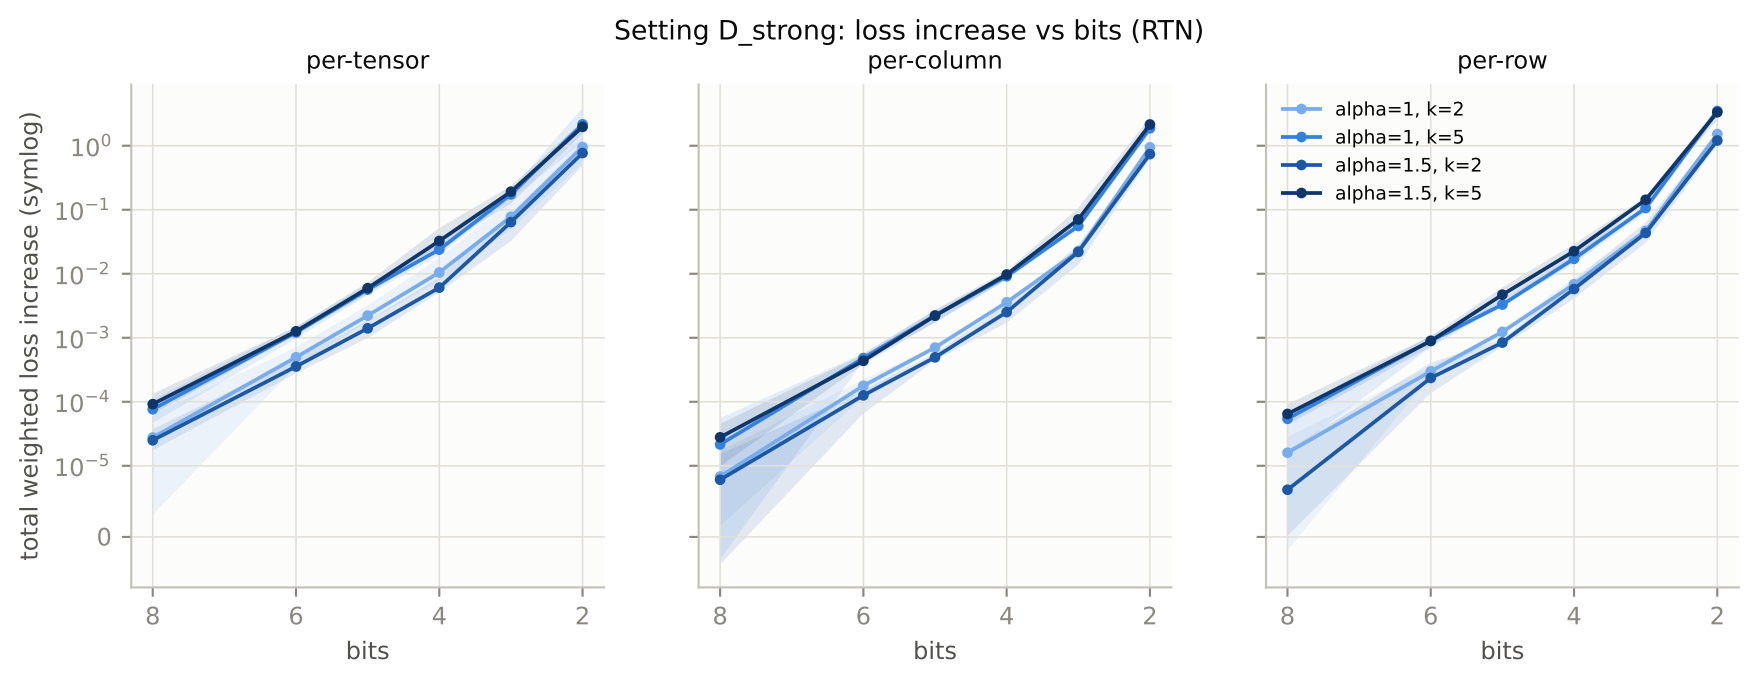
Setting D_strong: loss increase vs bits (RTN)
per-tensor
per-column
per-row
total weighted loss increase (symlog)
10^{0}
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
0
8
6
4
2
bits
alpha=1, k=2
alpha=1, k=5
alpha=1.5, k=2
alpha=1.5, k=5