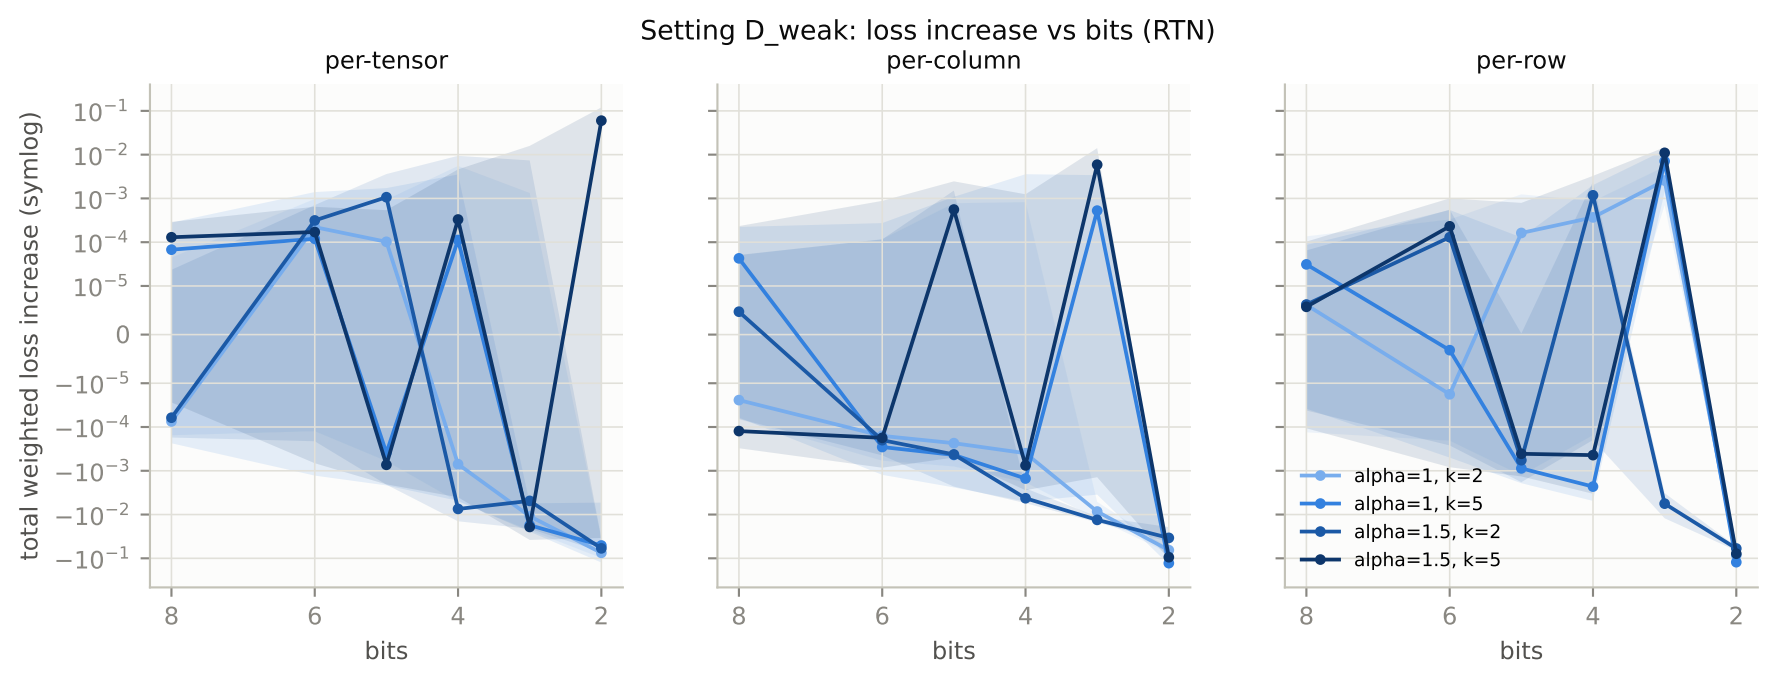
Setting D_weak: loss increase vs bits (RTN)
per-tensor
per-column
per-row
total weighted loss increase (symlog)
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
0
-10^{-5}
-10^{-4}
-10^{-3}
-10^{-2}
-10^{-1}
8
6
4
2
bits
alpha=1, k=2
alpha=1, k=5
alpha=1.5, k=2
alpha=1.5, k=5