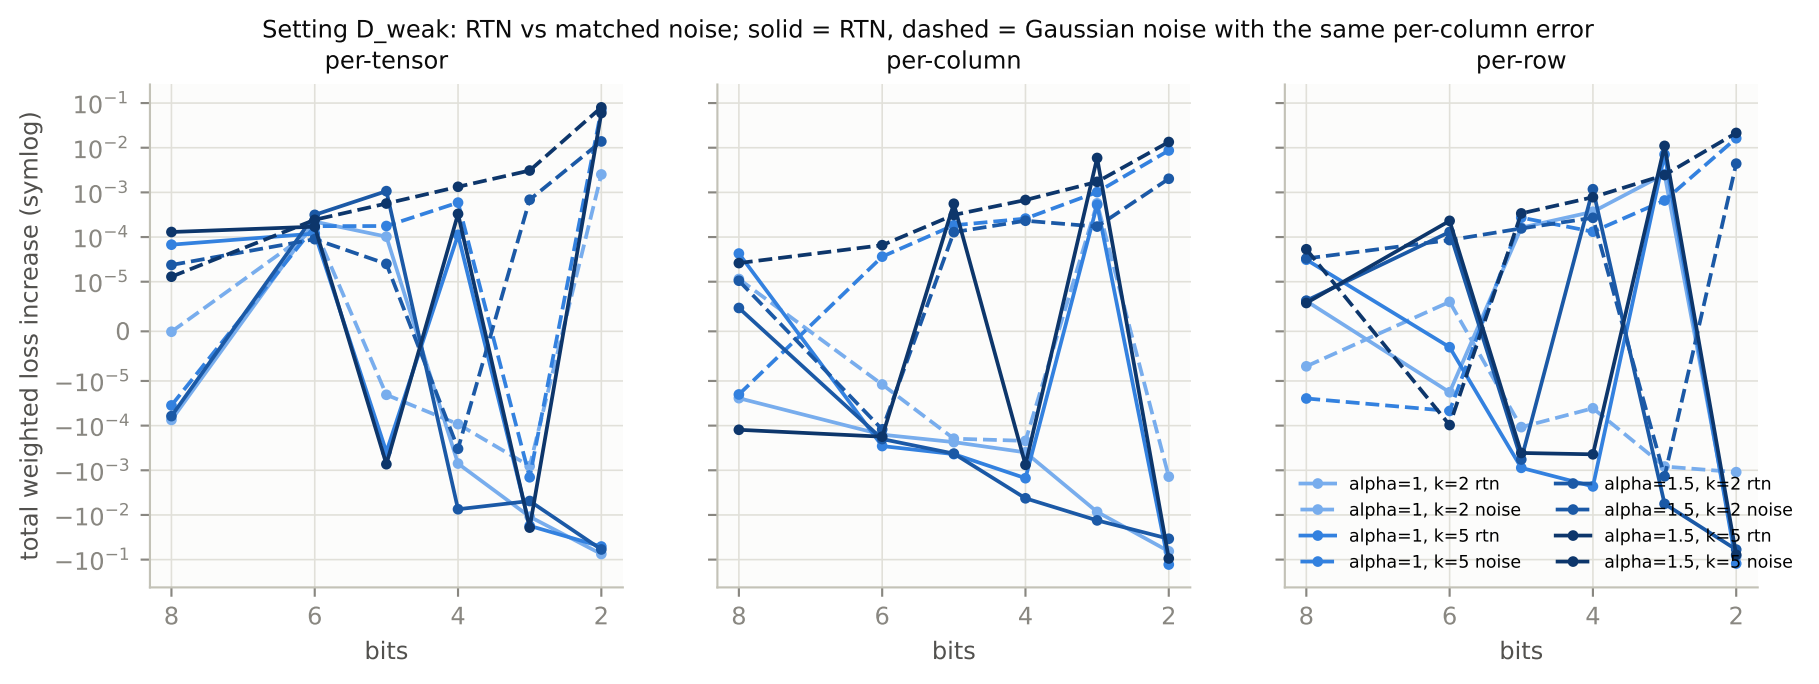
Setting D_weak: RTN vs matched noise; solid = RTN, dashed = Gaussian noise with the same per-column error
per-tensor
per-column
per-row
total weighted loss increase (symlog)
bits
alpha=1, k=2 rtn
alpha=1, k=2 noise
alpha=1, k=5 rtn
alpha=1, k=5 noise
alpha=1.5, k=2 rtn
alpha=1.5, k=2 noise
alpha=1.5, k=5 rtn
alpha=1.5, k=5 noise